

# Step 1: Data Import

<https://www.kaggle.com/code/ryuodan/asl-detection-walkthrough>

<https://www.kaggle.com/code/razinw/asl-alphabet-classification-with-cnn>

```
In [1]: import os
        from os import listdir
        from os.path import isfile, join
        import pandas as pd
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import accuracy_score
        from sklearn.metrics import confusion_matrix
        from sklearn.metrics import recall_score
        from sklearn.metrics import precision_score
        from sklearn.metrics import f1_score
        from sklearn.metrics import plot_roc_curve
        import random
        import matplotlib.image as mpimg
        import matplotlib.pyplot as plt
        import numpy as np
        import cv2
        import tensorflow as tf
        from tensorflow.keras import utils
        from tensorflow.keras.optimizers import Adam
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D,
        from tensorflow.keras import layers, models
        from sklearn import preprocessing
        from sklearn.metrics import confusion_matrix
        from scikeras.wrappers import KerasClassifier
        from sklearn.model_selection import cross_val_score
        import keras
        import keras_metrics
```

We iterate through each subfolder in the train folder, and for each subfolder, we add the paths of the image files to our list. Then we convert this to a dataframe.

```
In [2]: all_data = []
        train_folder = r'archive\asl_alphabet_train\asl_alphabet_train'

        for folder in os.listdir(train_folder):
            label_folder = os.path.join(train_folder, folder)
            onlyfiles = [{'label': folder, 'path': os.path.join(label_folder, f)} for f in os.
                          if os.path.isfile(os.path.join(label_folder, f))]
            all_data += onlyfiles


        data_df = pd.DataFrame(all_data)
```

## Step 2: Data Exploration

We perform exploratory analysis in order to view the distribution of our classes in our dataset. In addition, we provide previews of the original image and resized images to give a high-level overview of the types of inputs our models will be performing predictive analyses on.

```
In [3]: counts = data_df.groupby(by='label').size().to_frame(name='vals')

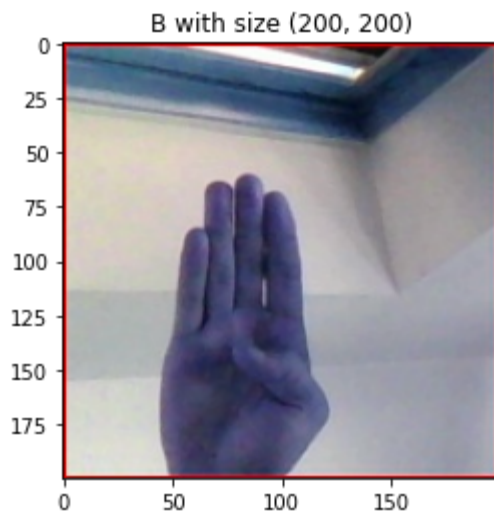
plt.figure(figsize=(20, 5))
plt.bar(counts.index, counts.vals)
plt.title('Counts of Classes')
plt.show()
```




```
In [4]: signs = data_df['label'].unique().tolist()

def show_image(label, signs, size):
    if label not in signs:
        print('Error: Unknown Label')
        return 0
    rows = data_df[data_df['label']==label]['path']
    rand_img = random.randint(a=0, b=len(rows))
    filepath = rows.iloc[rand_img]
    img = cv2.imread(filepath)
    img = cv2.resize(img, size)
    plt.figure()
    plt.title(label + " with size " + str(size))
    plt.imshow(img)
```


```
In [5]: original_size = 200, 200
show_image('B', signs, original_size)
```



```
In [6]: ideal_size = 32, 32  
show_image('B', signs, ideal_size)
```



```
In [7]: reduced_size = 20, 20  
show_image('B', signs, reduced_size)
```



## Step 3: Data Processing

We use the cv2 package to read an image from a filepath and resize it. We then convert the image to a numpy array of values between 0 and 255. This way, we can feed our data into our models.

```
In [8]: def encode_imgs(paths, size):  
        res = []  
        for path in paths:  
            img = cv2.imread(path)  
            img = cv2.resize(img, size)  
            res.append(img)  
        res = np.array(res)  
        res = res.astype('float32')/255.0  
        return res
```

## Step 4: Data Modeling

First, we establish a logistic regression model as a baseline, mainly to illustrate that this problem is solvable through machine learning. Logistic regression is relatively more simple than the neural networks we build, so we predict our neural networks will perform at least as well as this logistic regression model.

```
In [9]: def run_logr(data_df, label1, label2, size):
        binary_df = data_df.loc[(data_df['label'] == label1) | (data_df['label'] == label2)]
        encoded = encode_imgs(binary_df['path'], size)
        print('Data Encoded!\n')
        X_train, X_test, y_train, y_test = train_test_split(encoded, binary_df['label'])
        X_train = X_train.reshape(X_train.shape[0], X_train.shape[1] * X_train.shape[2])
        X_test = X_test.reshape(X_test.shape[0], X_test.shape[1] * X_test.shape[2] * X_test.shape[3])

        lr = LogisticRegression()
        lr.fit(X_train, y_train)
        print('Model fitted!\n')

        y_pred_test = lr.predict(X_test)
        print('Predictions made!\n')
        print('Confusion Matrix: \n')
        print(confusion_matrix(y_test, y_pred_test))
        test_acc = accuracy_score(y_test, y_pred_test)
        test_f1 = f1_score(y_test, y_pred_test, pos_label=label1)
        test_rec = recall_score(y_test, y_pred_test, pos_label=label1)
        test_pre = precision_score(y_test, y_pred_test, pos_label=label1)
        plot_roc_curve(lr, X_test, y_test)
        plt.show()

        print('Logistic Regression Metrics: ' + label1 + ' and ' + label2)
        print('Testing Accuracy: ' + str(test_acc))
        print('Testing F1 Score: ' + str(test_f1))
        print('Testing Recall: ' + str(test_rec))
        print('Testing Precision: ' + str(test_pre))
```

## Logistic Regression: Two "Similar" Signs

The letter A and the letter B have similar signs, so let's see how our Logistic Regression model performs.

```
In [10]: run_logr(data_df, 'A', 'B', ideal_size)
```

Data Encoded!

```
C:\Users\Justin\AppData\Local\Programs\Python\Python37\lib\site-packages\sklearn\linear_model\_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)


```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
C:\Users\Justin\AppData\Local\Programs\Python\Python37\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.
warnings.warn(msg, category=FutureWarning)
```

Model fitted!

Predictions made!

Confusion Matrix:

```
[[640  72]
 [ 76 712]]
```



Logistic Regression Metrics: A and B

Testing Accuracy: 0.9013333333333333

Testing F1 Score: 0.896358543417367

Testing Recall: 0.898876404494382

Testing Precision: 0.8938547486033519

## Logistic Regression: Two "Different" Signs

The letters C and G are vastly different signs, so let's see how our Logistic Regression model performs.

```
In [11]: run_logr(data_df, 'C', 'G', ideal_size)
```

Data Encoded!

```
C:\Users\Justin\AppData\Local\Programs\Python\Python37\lib\site-packages\sklearn\linear_model\_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)


```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
C:\Users\Justin\AppData\Local\Programs\Python\Python37\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.
warnings.warn(msg, category=FutureWarning)
```

Model fitted!

Predictions made!

Confusion Matrix:

```
[[712  0]
 [ 0 788]]
```



Logistic Regression Metrics: C and G

Testing Accuracy: 1.0

Testing F1 Score: 1.0

Testing Recall: 1.0

Testing Precision: 1.0

## Building Neural Networks

First, we encode the data to ideal size (32, 32) and the reduced size (20, 20).

```
In [12]: encoded_data = encode_imgs(data_df['path'], ideal_size)
reduced_data = encode_imgs(data_df['path'], reduced_size)
```

```
In [13]: le = preprocessing.LabelEncoder()
le.fit(data_df.label)
data_df['categorical_label'] = le.transform(data_df.label)
```

# NN 1

3 hidden layers (20, 40, 40 filters), input size (20, 20)

# NN 2

3 hidden layers (32, 64, 64 filters), input size (32, 32)

# NN 3

3 hidden layers (45, 90, 90 filters), input size (32, 32)

```
In [14]: model1 = models.Sequential()
model1.add(layers.Conv2D(20, (3, 3), activation='relu', input_shape=(20, 20, 3)))
model1.add(layers.MaxPooling2D((2, 2)))
model1.add(layers.Conv2D(40, (3, 3), activation='relu'))
model1.add(layers.MaxPooling2D((2, 2)))
model1.add(layers.Conv2D(40, (3, 3), activation='relu'))
model1.add(layers.Flatten())
model1.add(layers.Dense(40, activation='relu'))
model1.add(layers.Dense(29))

model1.summary()

model2 = models.Sequential()
model2.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.Flatten())
model2.add(layers.Dense(64, activation='relu'))
model2.add(layers.Dense(29))

model2.summary()

model3 = models.Sequential()
model3.add(layers.Conv2D(45, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model3.add(layers.MaxPooling2D((2, 2)))
model3.add(layers.Conv2D(90, (3, 3), activation='relu'))
model3.add(layers.MaxPooling2D((2, 2)))
model3.add(layers.Conv2D(90, (3, 3), activation='relu'))
model3.add(layers.Flatten())
model3.add(layers.Dense(90, activation='relu'))
model3.add(layers.Dense(29))

model3.summary()
```



Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 18, 18, 20)	560
max_pooling2d (MaxPooling2D)	(None, 9, 9, 20)	0
conv2d_1 (Conv2D)	(None, 7, 7, 40)	7240
max_pooling2d_1 (MaxPooling2D)	(None, 3, 3, 40)	0
conv2d_2 (Conv2D)	(None, 1, 1, 40)	14440
flatten (Flatten)	(None, 40)	0
dense (Dense)	(None, 40)	1640
dense_1 (Dense)	(None, 29)	1189

=====  
Total params: 25,069

Trainable params: 25,069

Non-trainable params: 0

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_2 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_4 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_5 (Conv2D)	(None, 4, 4, 64)	36928
flatten_1 (Flatten)	(None, 1024)	0
dense_2 (Dense)	(None, 64)	65600
dense_3 (Dense)	(None, 29)	1885

=====  
Total params: 123,805

Trainable params: 123,805

Non-trainable params: 0

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 30, 30, 45)	1260

max_pooling2d_4 (MaxPooling 2D)	(None, 15, 15, 45)	0
conv2d_7 (Conv2D)	(None, 13, 13, 90)	36540
max_pooling2d_5 (MaxPooling 2D)	(None, 6, 6, 90)	0
conv2d_8 (Conv2D)	(None, 4, 4, 90)	72990
flatten_2 (Flatten)	(None, 1440)	0
dense_4 (Dense)	(None, 90)	129690
dense_5 (Dense)	(None, 29)	2639

=====  
Total params: 243,119  
Trainable params: 243,119  
Non-trainable params: 0  
=====

## NN 4

1 hidden layer (32 filters), input size (32, 32)

```
In [15]: model4 = models.Sequential()  
model4.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))  
model4.add(layers.MaxPooling2D((2, 2)))  
model4.add(layers.Flatten())  
model4.add(layers.Dense(32, activation='relu'))  
model4.add(layers.Dense(29))  
  
model4.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_6 (MaxPooling 2D)	(None, 15, 15, 32)	0
flatten_3 (Flatten)	(None, 7200)	0
dense_6 (Dense)	(None, 32)	230432
dense_7 (Dense)	(None, 29)	957

=====  
Total params: 232,285  
Trainable params: 232,285  
Non-trainable params: 0  
=====

```
In [16]: def run_nn(model, optimizer, epochs, batch_size, data_df, data):
X_train, X_test, y_train, y_test = train_test_split(data, data_df['categorical_

model.compile(optimizer=optimizer,
               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True,
               metrics=['accuracy']))

history = model.fit(X_train, y_train, epochs=5, batch_size = 64, shuffle=True,
metrics = pd.DataFrame(history.history))

print('\n-----\n')
print(metrics)

y_pred_test = model.predict(X_test)
y_test_classes = np.argmax(y_pred_test,axis=1)

confusion = confusion_matrix(y_test, y_test_classes)
confusion_df = pd.DataFrame(confusion)
confusion_df.index = sorted(list(set(data_df['label'])))

acc_test = accuracy_score(y_test, y_test_classes)
f1_test = f1_score(y_test, y_test_classes, average='weighted')
prec_test = precision_score(y_test, y_test_classes, average='weighted')
rec_test = recall_score(y_test, y_test_classes, average='weighted')

print("\n-----\n")
print("Accuracy on testing set: ", acc_test)
print("F1 score on testing set: ", f1_test)
print("Precision on testing set: ", prec_test)
print("Recall on testing set: ", rec_test)

return confusion_df
```

```
In [17]: confusion_df1 = run_nn(model1, 'adam', 5, 64, data_df, reduced_data)
confusion_df1
```

```
Epoch 1/5
1020/1020 [=====] - 21s 20ms/step - loss: 2.2994 - accurac
y: 0.3095 - val_loss: 1.6304 - val_accuracy: 0.4850
Epoch 2/5
1020/1020 [=====] - 20s 19ms/step - loss: 1.3015 - accurac
y: 0.5826 - val_loss: 1.1617 - val_accuracy: 0.6150
Epoch 3/5
1020/1020 [=====] - 20s 19ms/step - loss: 0.9439 - accurac
y: 0.6885 - val_loss: 0.8558 - val_accuracy: 0.7206
Epoch 4/5
1020/1020 [=====] - 20s 19ms/step - loss: 0.7442 - accurac
y: 0.7521 - val_loss: 0.7209 - val_accuracy: 0.7572
Epoch 5/5
1020/1020 [=====] - 20s 19ms/step - loss: 0.6109 - accurac
y: 0.7926 - val_loss: 0.5775 - val_accuracy: 0.8051
```

```
-----

      loss  accuracy  val_loss  val_accuracy
0  2.299420  0.309487  1.630377      0.485011
1  1.301513  0.582559  1.161728      0.615034
2  0.943867  0.688460  0.855777      0.720552
3  0.744210  0.752077  0.720926      0.757241
4  0.610949  0.792613  0.577513      0.805149
```

```
-----

Accuracy on testing set: 0.8051494252873563
F1 score on testing set: 0.8051631258817623
Precision on testing set: 0.8176217205303256
Recall on testing set: 0.8051494252873563
```

Out[17]:

	0	1	2	3	4	5	6	7	8	9	...	19	20	21	22	23	24	2
<b>A</b>	555	30	0	4	57	5	0	0	3	0	...	4	0	0	0	0	3	
<b>B</b>	3	576	1	4	79	2	1	0	16	0	...	0	7	16	1	0	0	
<b>C</b>	5	2	689	8	1	4	0	0	0	0	...	0	1	0	0	0	0	
<b>D</b>	1	11	2	647	19	31	2	0	6	0	...	1	2	0	0	0	0	
<b>E</b>	24	40	0	4	607	25	7	0	11	0	...	1	0	0	0	0	0	
<b>F</b>	0	0	1	0	7	648	19	0	15	1	...	0	0	0	0	0	15	
<b>G</b>	0	0	0	0	0	9	639	44	7	19	...	0	0	0	0	0	0	
<b>H</b>	0	0	0	0	0	0	25	695	3	5	...	0	0	0	0	0	0	
<b>I</b>	0	0	0	2	5	4	19	0	664	20	...	0	0	9	1	1	1	
<b>J</b>	0	0	0	0	0	0	6	4	5	709	...	0	0	0	0	0	1	
<b>K</b>	0	3	0	1	0	0	0	0	49	0	...	0	3	24	13	0	0	
<b>L</b>	13	0	0	2	0	8	0	0	0	0	...	15	1	1	0	0	0	
<b>M</b>	10	0	0	0	2	1	0	0	7	8	...	1	0	1	8	4	5	
<b>N</b>	7	0	0	0	2	5	6	0	2	30	...	1	0	2	9	2	39	
<b>O</b>	4	0	10	5	2	17	3	0	1	0	...	5	1	8	9	14	24	1
<b>P</b>	0	0	0	0	0	2	5	5	0	4	...	0	0	0	0	0	0	
<b>Q</b>	0	0	3	0	0	0	1	0	0	0	...	0	0	0	0	0	5	
<b>R</b>	0	4	0	0	0	0	10	0	14	9	...	0	114	36	1	4	2	
<b>S</b>	4	1	1	0	1	15	0	0	0	0	...	34	2	4	3	3	26	
<b>T</b>	9	0	0	0	4	8	0	0	0	0	...	525	0	1	1	12	32	1
<b>U</b>	4	28	0	0	2	0	0	0	2	0	...	0	391	101	5	20	1	
<b>V</b>	1	2	0	2	0	5	3	0	19	0	...	1	22	518	71	8	7	
<b>W</b>	0	8	0	1	6	0	4	0	14	0	...	0	4	94	525	5	9	
<b>X</b>	0	0	0	0	0	15	0	0	14	4	...	12	21	31	2	501	52	1
<b>Y</b>	0	0	0	0	0	1	0	0	1	31	...	2	0	0	1	0	640	
<b>Z</b>	3	0	4	1	0	1	0	0	0	0	...	30	5	2	0	21	112	58
<b>del</b>	0	0	0	1	1	0	3	0	0	1	...	0	1	0	0	5	8	
<b>nothing</b>	0	0	0	0	0	0	1	0	0	0	...	0	0	0	0	0	0	
<b>space</b>	0	0	0	0	0	0	2	0	0	11	...	0	0	0	0	0	20	

29 rows × 29 columns

```
In [18]: confusion_df2 = run_nn(model2, 'adam', 5, 64, data_df, encoded_data)
          confusion_df2
```

```
Epoch 1/5
1020/1020 [=====] - 78s 76ms/step - loss: 1.4665 - accurac
y: 0.5528 - val_loss: 0.5729 - val_accuracy: 0.8127
Epoch 2/5
1020/1020 [=====] - 78s 76ms/step - loss: 0.3923 - accurac
y: 0.8691 - val_loss: 0.3030 - val_accuracy: 0.8900
Epoch 3/5
1020/1020 [=====] - 77s 76ms/step - loss: 0.2001 - accurac
y: 0.9331 - val_loss: 0.1524 - val_accuracy: 0.9492
Epoch 4/5
1020/1020 [=====] - 77s 76ms/step - loss: 0.1225 - accurac
y: 0.9596 - val_loss: 0.1301 - val_accuracy: 0.9551
Epoch 5/5
1020/1020 [=====] - 78s 77ms/step - loss: 0.0859 - accurac
y: 0.9722 - val_loss: 0.0921 - val_accuracy: 0.9696
```

-----

	loss	accuracy	val_loss	val_accuracy
0	1.466454	0.552751	0.572889	0.812690
1	0.392285	0.869149	0.303043	0.889977
2	0.200150	0.933073	0.152425	0.949195
3	0.122524	0.959602	0.130149	0.955126
4	0.085860	0.972184	0.092085	0.969609

-----

```
Accuracy on testing set: 0.9696091954022988
F1 score on testing set: 0.969329590654926
Precision on testing set: 0.9712433772171185
Recall on testing set: 0.9696091954022988
```

Out[18]:

	0	1	2	3	4	5	6	7	8	9	...	19	20	21	22	23	24	2
<b>A</b>	715	0	0	1	0	0	0	0	0	0	...	2	0	0	0	0	0	
<b>B</b>	0	714	0	7	5	0	0	0	1	0	...	0	0	0	0	0	0	
<b>C</b>	0	0	727	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>D</b>	0	0	0	753	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>E</b>	8	2	0	7	734	6	0	0	0	0	...	0	0	0	1	0	0	
<b>F</b>	0	0	0	0	0	724	0	0	0	0	...	0	0	0	0	0	0	
<b>G</b>	0	0	0	0	0	7	711	0	1	0	...	0	0	0	0	0	0	4
<b>H</b>	0	0	0	0	0	0	7	732	0	0	...	0	0	0	0	0	0	0
<b>I</b>	0	0	0	0	0	3	0	0	744	0	...	0	0	0	0	0	0	1
<b>J</b>	0	0	0	0	0	0	2	0	1	728	...	0	0	0	0	0	0	5
<b>K</b>	0	0	0	1	0	0	0	0	0	0	...	0	0	5	4	0	0	
<b>L</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>M</b>	5	0	0	0	1	0	0	0	0	1	...	0	0	0	8	0	0	
<b>N</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>O</b>	0	0	0	2	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>P</b>	0	0	0	0	0	0	0	0	0	3	...	0	0	0	0	0	0	
<b>Q</b>	0	0	0	0	0	0	0	0	0	0	...	1	0	0	0	0	0	
<b>R</b>	0	0	0	0	0	0	0	0	0	0	...	0	13	3	2	2	0	
<b>S</b>	5	0	0	0	1	2	0	0	0	0	...	52	0	0	0	11	16	
<b>T</b>	0	0	0	0	0	2	0	0	0	0	...	718	0	0	0	2	1	1
<b>U</b>	0	0	0	1	0	0	0	0	0	0	...	0	696	2	4	7	0	
<b>V</b>	0	0	0	0	0	0	0	0	4	0	...	0	0	627	100	3	1	
<b>W</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	3	748	1	0	
<b>X</b>	0	0	0	0	0	1	0	0	0	0	...	1	0	0	6	693	6	1
<b>Y</b>	0	0	0	0	0	1	0	0	0	0	...	0	0	0	0	2	706	
<b>Z</b>	0	0	0	0	0	0	0	0	0	0	...	0	2	0	0	0	0	80
<b>del</b>	0	0	0	0	1	0	1	0	0	0	...	0	0	0	0	0	0	
<b>nothing</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>space</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	1

29 rows × 29 columns

```
In [19]: confusion_df3 = run_nn(model3, 'adam', 5, 64, data_df, encoded_data)
          confusion_df3
```

```
Epoch 1/5
1020/1020 [=====] - 118s 114ms/step - loss: 1.3295 - accur
acy: 0.5916 - val_loss: 0.4815 - val_accuracy: 0.8393
Epoch 2/5
1020/1020 [=====] - 116s 114ms/step - loss: 0.3239 - accur
acy: 0.8930 - val_loss: 0.2108 - val_accuracy: 0.9303
Epoch 3/5
1020/1020 [=====] - 115s 112ms/step - loss: 0.1550 - accur
acy: 0.9492 - val_loss: 0.0984 - val_accuracy: 0.9689
Epoch 4/5
1020/1020 [=====] - 115s 113ms/step - loss: 0.1004 - accur
acy: 0.9671 - val_loss: 0.1049 - val_accuracy: 0.9626
Epoch 5/5
1020/1020 [=====] - 119s 117ms/step - loss: 0.0701 - accur
acy: 0.9779 - val_loss: 0.0802 - val_accuracy: 0.9720
```

-----

	loss	accuracy	val_loss	val_accuracy
0	1.329503	0.591586	0.481489	0.839264
1	0.323881	0.892981	0.210805	0.930299
2	0.155001	0.949241	0.098398	0.968874
3	0.100413	0.967111	0.104922	0.962575
4	0.070145	0.977870	0.080207	0.971954

-----

```
Accuracy on testing set: 0.9719540229885058
F1 score on testing set: 0.9719503401826124
Precision on testing set: 0.973931597908162
Recall on testing set: 0.9719540229885058
```



Out[19]:

	0	1	2	3	4	5	6	7	8	9	...	19	20	21	22	23	24	2
<b>A</b>	719	2	0	0	3	0	0	0	0	0	...	2	0	0	0	0	0	
<b>B</b>	0	736	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>C</b>	0	0	727	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>D</b>	0	0	1	738	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>E</b>	8	33	1	0	702	2	0	0	0	0	...	2	0	0	0	0	0	
<b>F</b>	0	0	0	0	0	723	0	0	0	0	...	0	0	0	0	0	0	
<b>G</b>	0	0	0	0	0	0	721	2	0	0	...	0	0	0	0	0	0	
<b>H</b>	0	0	0	0	0	0	0	739	0	0	...	0	0	0	0	0	0	
<b>I</b>	0	4	0	2	2	0	0	4	724	0	...	0	0	0	9	0	0	
<b>J</b>	0	0	0	0	0	0	0	0	0	723	...	0	0	0	0	0	13	
<b>K</b>	0	8	0	0	0	0	0	0	4	0	...	0	8	6	16	0	0	
<b>L</b>	0	0	0	0	0	0	0	0	0	0	...	1	0	0	0	0	0	
<b>M</b>	3	0	0	0	0	0	0	0	0	0	...	0	0	0	0	1	0	
<b>N</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	1	0	0	
<b>O</b>	2	0	3	0	0	0	0	0	0	0	...	5	0	0	0	0	0	
<b>P</b>	0	0	0	0	0	0	0	1	0	0	...	1	0	0	0	0	2	
<b>Q</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>R</b>	0	3	0	0	0	0	0	0	0	0	...	0	108	3	1	0	0	
<b>S</b>	0	0	0	0	0	0	0	0	3	0	...	27	0	0	0	15	1	
<b>T</b>	0	0	0	0	0	0	0	0	0	0	...	729	0	0	0	1	9	
<b>U</b>	0	1	0	0	0	0	0	0	0	0	...	0	710	3	6	15	0	
<b>V</b>	0	5	0	0	0	0	0	0	0	0	...	0	24	619	94	10	0	
<b>W</b>	0	3	0	0	0	0	0	0	0	0	...	0	0	0	752	0	0	
<b>X</b>	2	0	0	0	0	0	0	0	0	0	...	0	6	0	0	715	1	
<b>Y</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	710	
<b>Z</b>	0	0	1	0	0	0	0	0	0	0	...	1	0	0	0	0	0	80
<b>del</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>nothing</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	
<b>space</b>	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	

29 rows × 29 columns

```
In [20]: confusion_df4 = run_nn(model4, 'adam', 5, 64, data_df, encoded_data)
          confusion_df4
```

```
Epoch 1/5
1020/1020 [=====] - 46s 44ms/step - loss: 2.7771 - accurac
y: 0.1560 - val_loss: 2.2759 - val_accuracy: 0.2635
Epoch 2/5
1020/1020 [=====] - 45s 44ms/step - loss: 2.0881 - accurac
y: 0.3065 - val_loss: 1.9190 - val_accuracy: 0.3652
Epoch 3/5
1020/1020 [=====] - 45s 44ms/step - loss: 1.8171 - accurac
y: 0.3849 - val_loss: 1.7247 - val_accuracy: 0.4143
Epoch 4/5
1020/1020 [=====] - 45s 44ms/step - loss: 1.6346 - accurac
y: 0.4356 - val_loss: 1.5727 - val_accuracy: 0.4472
Epoch 5/5
1020/1020 [=====] - 45s 44ms/step - loss: 1.5049 - accurac
y: 0.4741 - val_loss: 1.4423 - val_accuracy: 0.4953
```

-----

	loss	accuracy	val_loss	val_accuracy
0	2.777118	0.155954	2.275889	0.263494
1	2.088121	0.306529	1.918957	0.365195
2	1.817104	0.384889	1.724670	0.414253
3	1.634560	0.435556	1.572677	0.447172
4	1.504934	0.474100	1.442312	0.495264

-----

```
Accuracy on testing set: 0.49526436781609195
F1 score on testing set: 0.48700689599581193
Precision on testing set: 0.49170432149624177
Recall on testing set: 0.49526436781609195
```

Out[20]:

	0	1	2	3	4	5	6	7	8	9	...	19	20	21	22	23	24	2
<b>A</b>	364	22	6	23	77	146	1	0	65	1	...	0	0	0	0	0	0	
<b>B</b>	52	265	0	49	3	256	0	0	19	5	...	0	0	0	2	0	0	
<b>C</b>	5	0	545	24	33	0	0	0	42	29	...	0	0	0	0	0	0	
<b>D</b>	29	68	20	370	20	114	0	0	97	0	...	2	0	0	0	0	0	
<b>E</b>	217	38	25	53	218	80	22	0	50	2	...	4	0	0	0	0	0	
<b>F</b>	71	87	16	24	70	359	0	0	53	0	...	0	0	0	0	0	0	
<b>G</b>	4	0	0	0	0	0	506	130	0	76	...	0	0	0	0	0	0	
<b>H</b>	0	0	0	0	0	0	257	376	0	68	...	0	0	0	0	0	0	
<b>I</b>	12	5	26	48	48	4	0	0	354	15	...	6	0	0	0	0	0	
<b>J</b>	0	0	11	0	0	0	21	103	6	435	...	0	0	0	0	0	0	
<b>K</b>	0	40	0	8	0	0	0	0	26	4	...	5	0	3	22	0	0	
<b>L</b>	2	4	25	1	1	0	0	0	96	0	...	1	0	0	0	0	0	
<b>M</b>	0	0	24	0	0	0	0	1	35	14	...	22	0	0	0	1	0	
<b>N</b>	0	0	13	0	0	0	7	23	10	55	...	12	0	0	0	0	1	
<b>O</b>	0	2	8	1	0	0	0	0	66	3	...	41	1	0	0	6	7	
<b>P</b>	0	0	0	0	0	0	0	32	0	23	...	0	0	0	0	0	0	
<b>Q</b>	0	0	1	0	0	0	0	0	0	3	...	3	0	0	0	0	0	
<b>R</b>	0	0	0	0	0	0	0	0	6	2	...	31	193	71	24	44	14	
<b>S</b>	0	1	0	0	0	0	0	0	5	21	...	33	10	1	4	39	104	6
<b>T</b>	1	0	0	0	0	0	0	0	37	0	...	239	11	0	0	48	12	15
<b>U</b>	0	0	0	0	0	0	0	0	3	7	...	23	291	66	65	66	38	
<b>V</b>	0	0	0	0	0	0	0	0	0	15	...	0	184	124	314	9	12	
<b>W</b>	0	0	0	0	0	0	0	0	0	0	...	1	109	110	438	7	5	
<b>X</b>	0	0	0	0	0	0	0	0	3	5	...	20	69	4	7	113	192	13
<b>Y</b>	0	0	0	0	0	0	0	1	1	18	...	20	31	1	10	42	323	3
<b>Z</b>	0	0	0	0	0	0	0	0	7	0	...	113	35	0	0	68	76	24
<b>del</b>	0	0	0	0	0	0	0	2	0	10	...	3	0	0	0	0	0	
<b>nothing</b>	0	0	0	0	0	0	0	0	0	5	...	0	0	0	0	0	0	
<b>space</b>	0	0	0	0	0	0	0	0	1	2	...	0	4	4	4	1	54	

29 rows × 29 columns

```
In [27]: def run_kfold_nn(model, epochs, batch_size, encoded_data, data_df, cv):  
         kc = KerasClassifier(model = model, epochs = epochs, batch_size = batch_size, v  
         cvs = cross_val_score(kc, encoded_data, data_df['label'], cv = cv)  
         return cvs
```

```
In [28]: cvs1 = run_kfold_nn(model1, 1, 64, reduced_data, data_df, 3)  
         cvs2 = run_kfold_nn(model2, 1, 64, encoded_data, data_df, 3)  
         cvs3 = run_kfold_nn(model3, 1, 64, encoded_data, data_df, 3)  
         cvs4 = run_kfold_nn(model4, 1, 64, encoded_data, data_df, 3)
```

```
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmp4_n1kb_z\assets
907/907 [=====] - 18s 19ms/step - loss: 0.4690 - accuracy: 0.8403
454/454 [=====] - 3s 7ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmpfh8i06jc\assets
907/907 [=====] - 18s 19ms/step - loss: 0.3891 - accuracy: 0.8679
454/454 [=====] - 3s 7ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmp2pd19ms5\assets
907/907 [=====] - 18s 19ms/step - loss: 0.5554 - accuracy: 0.8116
454/454 [=====] - 3s 7ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmpdwaohslq\assets
907/907 [=====] - 52s 57ms/step - loss: 0.0579 - accuracy: 0.9819
454/454 [=====] - 7s 15ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmpjuvrd701\assets
907/907 [=====] - 46s 50ms/step - loss: 0.0619 - accuracy: 0.9798
454/454 [=====] - 8s 18ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmp9lt4ik9l\assets
907/907 [=====] - 50s 54ms/step - loss: 0.0829 - accuracy: 0.9733
454/454 [=====] - 7s 15ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmp4y2pw4qo\assets
907/907 [=====] - 75s 82ms/step - loss: 0.0458 - accuracy: 0.9860
454/454 [=====] - 11s 23ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmp7996ugnq\assets
907/907 [=====] - 70s 77ms/step - loss: 0.0465 - accuracy: 0.9855
454/454 [=====] - 11s 22ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmp9csomr7o\assets
907/907 [=====] - 72s 78ms/step - loss: 0.0666 - accuracy: 0.9797
454/454 [=====] - 10s 23ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmp3254mn2u\assets
907/907 [=====] - 24s 26ms/step - loss: 1.3396 - accuracy: 0.5298
454/454 [=====] - 4s 9ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmp_n3gl1aqe\assets
907/907 [=====] - 24s 26ms/step - loss: 1.2352 - accuracy: 0.5649
454/454 [=====] - 4s 10ms/step
INFO:tensorflow:Assets written to: C:\Users\Justin\AppData\Local\Temp\tmpbotsmpcw\assets
907/907 [=====] - 27s 29ms/step - loss: 1.5327 - accuracy: 0.4608
```

454/454 [=====] - 5s 10ms/step

```
In [29]: print(cvs1)
         print(cvs2)
         print(cvs3)
         print(cvs4)
```

```
[0.74272414 0.64634483 0.80265517]
[0.94337931 0.84306897 0.96455172]
[0.88506897 0.92241379 0.95137931]
[0.42844828 0.32713793 0.52306897]
```